NLP Programming Tutorial 2 - Bigram Language Models

Graham Neubig
Nara Institute of Science and Technology (NAIST)

Review:
Calculating Sentence Probabilities

- We want the probability of

$$W = \text{speech recognition system}$$

- Represent this mathematically as:

$P(|W| = 3, w_1=\text{"speech"}, w_2=\text{"recognition"}, w_3=\text{"system"}) =$

$P(w_1=\text{"speech"} \mid w_0 = \text{"<s>"})$

$* P(w_2=\text{"recognition"} \mid w_0 = \text{"<s>"}, w_1=\text{"speech"})$

$* P(w_3=\text{"system"} \mid w_0 = \text{"<s>"}, w_1=\text{"speech"}, w_2=\text{"recognition"})$

$* P(w_4=\text{"</s>"} \mid w_0 = \text{"<s>"}, w_1=\text{"speech"}, w_2=\text{"recognition"}, w_3=\text{"system"})$

NOTE:
sentence start <s> and end </s> symbol

NOTE:
$P(w_0 = \text{<s>}) = 1$

2

Incremental Computation

- Previous equation can be written:

$$P(W) = \prod_{i=1}^{|W|+1} P(w_i | w_0 \dots w_{i-1})$$

- Unigram model ignored context:

$$P(w_i | w_0 \dots w_{i-1}) \approx P(w_i)$$

Unigram Models Ignore Word Order!

- Ignoring context, probabilities are the same:

P_{uni}(w=speech recognition system) =
  P(w=speech) * P(w=recognition) * P(w=system) * P(w=</s>)

=

P_{uni}(w=system recognition speech) =
  P(w=speech) * P(w=recognition) * P(w=system) * P(w=</s>)

4

Unigram Models Ignore Agreement!

- Good sentences (words agree):

P_{uni}(w=i am) =
  P(w=i) * P(w=am) * P(w=</s>)

P_{uni}(w=we are) =
  P(w=we) * P(w=are) * P(w=</s>)

- Bad sentences (words don't agree)

P_{uni}(w=we am) =
  P(w=we) * P(w=am) * P(w=</s>)

P_{uni}(w=i are) =
  P(w=i) * P(w=are) * P(w=</s>)

But no penalty because probabilities are independent!

Solution: Add More Context!

- Unigram model ignored context:

$$P(w_i|w_0\ldots w_{i-1}) \approx P(w_i)$$

- Bigram model adds one word of context

$$P(w_i|w_0\ldots w_{i-1}) \approx P(w_i|w_{i-1})$$

- Trigram model adds two words of context

$$P(w_i|w_0\ldots w_{i-1}) \approx P(w_i|w_{i-2}w_{i-1})$$

- Four-gram, five-gram, six-gram, etc...

6

Maximum Likelihood Estimation of n-gram Probabilities

- Calculate counts of n word and n-1 word strings

$$P(w_i | w_{i-n+1} \ldots w_{i-1}) = \frac{c(w_{i-n+1} \ldots w_i)}{c(w_{i-n+1} \ldots w_{i-1})}$$

i live in osaka . </s>
i am a graduate student . </s>
my school is in nara . </s>

n=2 →

P(osaka | in) = c(in osaka)/c(in) = 1 / 2 = 0.5
P(nara | in) = c(in nara)/c(in) = 1 / 2 = 0.5

7

Still Problems of Sparsity

- When n-gram frequency is 0, probability is 0

P(osaka | in) = c(i osaka)/c(in) = 1 / 2 = 0.5

P(nara | in) = c(i nara)/c(in) = 1 / 2 = 0.5

P(school | in) = c(in school)/c(in) = 0 / 2 = **0**!!

- Like unigram model, we can use linear interpolation

Bigram:
$$P(w_i|w_{i-1}) = \lambda_2 P_{ML}(w_i|w_{i-1}) + (1-\lambda_2) P(w_i)$$

Unigram:
$$P(w_i) = \lambda_1 P_{ML}(w_i) + (1-\lambda_1)\frac{1}{N}$$

Choosing Values of λ: Grid Search

- One method to choose λ_2, λ_1: try many values

$$\lambda_2 = 0.95, \lambda_1 = 0.95$$
$$\lambda_2 = 0.95, \lambda_1 = 0.90$$
$$\lambda_2 = 0.95, \lambda_1 = 0.85$$
$$\ldots$$
$$\lambda_2 = 0.95, \lambda_1 = 0.05$$
$$\lambda_2 = 0.90, \lambda_1 = 0.95$$
$$\lambda_2 = 0.90, \lambda_1 = 0.90$$
$$\ldots$$
$$\lambda_2 = 0.05, \lambda_1 = 0.10$$
$$\lambda_2 = 0.05, \lambda_1 = 0.05$$

Problems:

Too many options
→ Choosing takes time!

Using same λ for all n-grams
→ There is a smarter way!

9

Context Dependent Smoothing

High frequency word: "Tokyo"

c(Tokyo city) = 40
c(Tokyo is) = 35
c(Tokyo was) = 24
c(Tokyo tower) = 15
c(Tokyo port) = 10

…

Most 2-grams already exist
→ Large λ is better!

Low frequency word: "Tottori"

c(Tottori is) = 2
c(Tottori city) = 1
c(Tottori was) = 0

Many 2-grams will be missing
→ Small λ is better!

- Make the interpolation depend on the context

$$P(w_i|w_{i-1}) = \lambda_{w_{i-1}} P_{ML}(w_i|w_{i-1}) + (1 - \lambda_{w_{i-1}}) P(w_i)$$

10

Witten-Bell Smoothing

- One of the many ways to choose $\lambda_{w_{i-1}}$

$$\lambda_{w_{i-1}} = 1 - \frac{u(w_{i-1})}{u(w_{i-1}) + c(w_{i-1})}$$

$u(w_{i-1})$ = number of <u>unique words</u> after w_{i-1}

- For example:

c(Tottori is) = 2   c(Tottori city) = 1
c(Tottori) = 3   u(Tottori) = 2

$$\lambda_{Tottori} = 1 - \frac{2}{2+3} = 0.6$$

c(Tokyo city) = 40 c(Tokyo is) = 35 ...
c(Tokyo) = 270   u(Tokyo) = 30

$$\lambda_{Tokyo} = 1 - \frac{30}{30+270} = 0.9$$

Programming Techniques

Inserting into Arrays

- To calculate n-grams easily, you may want to:

my_words = ["this", "is", "a", "pen"]

↓

my_words = ["<s>", "this", "is", "a", "pen", "</s>"]

- This can be done with:

my_words.**append**("</s>")   # Add to the end
my_words.**insert**(0, "<s>")   # Add to the beginning

13

Removing from Arrays

- Given an n-gram with w_{i-n+1} … w_i, we may want the context w_{i-n+1} … w_{i-1}

- This can be done with:

```
my_ngram = "tokyo tower"
my_words = my_ngram.split(" ")    # Change into ["tokyo", "tower"]
my_words.pop()                    # Remove the last element ("tower")
my_context = " ".join(my_words)   # Join the array back together
print my_context
```

Exercise

Exercise

- Write two programs

  - train-bigram: Creates a bigram model

  - test-bigram: Reads a bigram model and calculates entropy on the test set

- Test train-bigram on test/02-train-input.txt

- Train the model on data/wiki-en-train.word

- Calculate entropy on data/wiki-en-test.word (if linear interpolation, test different values of λ_2)

- Challenge:

  - Use Witten-Bell smoothing (Linear interpolation is easier)

  - Create a program that works with *any* n (not just bi-gram)

16

train-bigram (Linear Interpolation)

create **map** *counts, context_counts*

for each *line* **in** the *training_file*
  split *line* into an array of *words*
  append "</s>" to the end and "<s>" to the beginning of *words*
  for each i **in** 1 **to** length(*words*)-1 # Note: starting at 1, **after** <s>
    counts["w_{i-1} w_i"] += 1   # Add bigram and bigram context

    context_counts["w_{i-1}"] += 1

    counts["w_i"] += 1   # Add unigram and unigram context

    context_counts[""] += 1

open the *model_file* for writing
for each *ngram*, *count* **in** *counts*
  split *ngram* into an array of *words*   # "w_{i-1} w_i" → {"w_{i-1}", "w_i"}
  remove the last element of *words*   # {"w_{i-1}", "w_i"} → {"w_{i-1}"}
  join *words* into *context*   # {"w_{i-1}"} → "w_{i-1}"
  probability = counts[ngram]/context_counts[context]
  print *ngram*, *probability* **to** *model_file*

17

test-bigram (Linear Interpolation)

λ_1 = ???, λ_2 = ???, V = 1000000, W = 0, H = 0

load model into *probs*

for each *line* **in** *test_file*
  split *line* into an array of *words*
  append "</s>" to the end and "<s>" to the beginning of *words*
  for each i **in** 1 **to** length(*words*)-1   # Note: starting at 1, **after** <s>
    P1 = λ_1 *probs*["w_i"] + (1 − λ_1) / V   # Smoothed unigram probability
    P2 = λ_2 *probs*["w_{i-1} w_i"] + (1 − λ_2) * P1 # Smoothed bigram probability
    H += -$\log_2$(P2)
    W += 1

print "entropy = "+*H/W*

18

Thank You!